

## WHITEPAPER

### Designing a Low-Latency Order Matching System for Real-Time Electronic Trading

**Author:** Ravuri Varshitha

**Date:** December 2025

**Version:** 1.0

---

#### Executive Summary

Electronic trading systems execute millions of buy/sell orders across global markets every second. In this environment, **latency is a competitive advantage**: faster systems execute more trades, capture better prices, and reduce risk exposure.

This whitepaper presents the architecture, components, data flow, and performance optimizations behind a **low-latency order matching system**, similar to the systems used by trading desks at firms such as Goldman Sachs, JP Morgan, and other global exchanges.

The design focuses on:

- **In-memory matching engine**
- **Price-time priority order book**
- **High-throughput event pipelines**
- **Asynchronous persistence**
- **Microsecond-level optimizations**

This paper is intended to demonstrate both **system design expertise** and **financial engineering awareness**.

---

#### 1. Problem Statement

Modern electronic trading requires:

- Handling **hundreds of thousands of orders per second**
- Executing trades with **microsecond to millisecond latency**
- Ensuring **consistency, fairness**, and **deterministic matching**
- Minimizing risk through **pre-trade risk checks**
- Broadcasting real-time **market data updates**

Traditional database-centric architectures cannot meet these performance requirements.

**Goal:**

Design a **low-latency order matching system** with:

- In-memory order book
  - Deterministic matching rules
  - Fault tolerance
  - Real-time market data dissemination
  - Asynchronous persistence for audit/compliance
- 

## 2. System Requirements

### 2.1 Functional Requirements

- Accept **BUY** and **SELL** orders
- Maintain a **real-time order book**
- Match orders using **price-time priority**
- Support order types (limit, market, cancel)
- Send **execution confirmations** to clients
- Publish **market data** (trades, quotes)

### 2.2 Non-Functional Requirements

#### Requirement Target

Latency      **< 1 millisecond** end-to-end

Throughput    **100K+ orders per second**

Availability    **99.99%**

Consistency   Deterministic matching

Durability     Asynchronous persistence

---

## 3. High-Level Architecture Overview

A low-latency trading system typically includes these components:

1. **Client Gateway**
2. **Order Normalizer & Validator**
3. **Risk Engine**
4. **Matching Engine** (core)
5. **In-Memory Order Book**
6. **Market Data Publisher**
7. **Event Logger → Asynchronous DB Writer**

### **Simplified Architecture Flow**

[Clients]

|

v

[Gateway] -> [Risk Checks] -> [Matching Engine] -> [Market Data Publisher]

|

v

[Event Logger → DB Writer → Database]

---

## **4. Component Breakdown**

### **4.1 Gateway**

- Accepts orders via FIX/WebSocket/REST
- Performs basic field validation
- Low-latency kernel-bypass networking (optional)

### **4.2 Normalizer**

- Converts all orders into an internal canonical format
- Ensures deterministic processing

### **4.3 Risk Engine**

Performs **pre-trade risk checks** such as:

- Margin availability
- Position limits

- Fat-finger checks
- Credit exposure

Crucially:

**Risk checks must also be near-instant — in-memory or cache-based.**

#### 4.4 Matching Engine (Core)

This component:

- Maintains order book in memory
- Inserts and matches orders
- Enforces **price-time priority**
- Generates trade events

Matching logic must complete in **microseconds**, so:

- NO database calls
- NO locking (prefer lock-free structures)
- NO GC interruptions
- **Why Matching Engines Are Often Single-Threaded**
- Although counterintuitive, many high-performance matching engines are:
- Single-threaded
- CPU-pinned
- Lock-free
- **Reason:**
- Locks introduce unpredictability
- Deterministic execution is critical
- Modern CPUs are extremely fast per core

#### 4.5 In-Memory Order Book

Two balanced priority queues:

BID side → max-heap (highest price first)

ASK side → min-heap (lowest price first)

If best bid  $\geq$  best ask  $\rightarrow$  match occurs.

## 4.6 Market Data Publisher

Publishes:

- Trades
- Top-of-book quotes
- Order book updates

Sent via:

- Kafka streams
- Redis pub/sub
- WebSocket servers

## 4.7 Persistent Storage (Async)

Trading systems can never block on DB writes.

Instead:

The matching engine emits **events** to an event log, e.g.:

- Kafka
- NATS
- Chronicle Queue

A separate writer service stores them in:

- PostgreSQL
- Cassandra
- Time-series DB

This ensures **ultra-low latency** in the critical path.

---

## 5. Data Flow

1. Client sends order
2. Gateway performs validation
3. Risk engine applies checks
4. Matching engine updates in-memory order book

5. Order is matched or stored
6. Execution event is generated
7. Market data publisher broadcasts update
8. Event logger saves event asynchronously
9. DB writer persists to database

No blocking database interaction until step 8.

---

## **6. Matching Engine Logic**

### **Price-Time Priority Rules**

1. Better price always wins
2. If prices are equal → earlier timestamp wins

### **Limit Order Flow**

If BUY\_LIMIT:

    check best ASK price

    if ASK  $\leq$  BUY:

        trade occurs

    else:

        add to BID book

### **Market Order Flow**

Match against best available prices

until quantity fulfilled or book exhausted

### **Cancel Order Flow**

- Remove from the corresponding queue
- Update book snapshot

### **Handling Partial Fills**

Orders may be partially matched; the remainder stays active.

---

## **7. Performance Optimizations**

To achieve sub-millisecond latency:

### 7.1 Hot Path Optimization

- Keep entire order book **in RAM**
- Avoid locking → use lock-free queues
- Avoid syscalls → kernel bypass (DPDK)
- Pre-allocate memory → avoid GC stalls
- Use CPU pinning
- NUMA-aware placement

### Latency Optimization Techniques

#### Hot Path vs Cold Path

- **Hot Path:**
  - Everything required to execute a trade
  - Must be extremely fast
- **Cold Path:**
  - Logging
  - Monitoring
  - Persistence
  - Analytics
- Key rule:
  - **Nothing unnecessary runs on the hot path**

#### --Memory Optimization

- Avoid object creation
- Reuse data structures
- Pre-allocate buffers
- Maintain cache locality

Garbage collection pauses are avoided by design.

#### Concurrency Strategy

- Minimal locking

- Event queues
- Single writer principle
- Carefully controlled threading

Concurrency bugs are more dangerous than small delay.

#### -- I/O Optimization

- No database calls in execution path
- Use append-only logs
- Batch disk writes
- Persist asynchronously

## 7.2 Asynchronous Persistence

Do NOT wait for DB writes.

Publish trade/order events to a log stream.

## 7.3 Network Optimizations

- TCP\_NODELAY
- Busy-polling sockets
- Zero-copy I/O

## 7.4 Scaling by Instrument

Each trading symbol (AAPL, NIFTY, BTC) can run on a **separate matching engine instance**, scaling horizontally.

---

## 8. Failure Handling & Fault Tolerance

### 8.1 Event Replay

If a matching engine crashes:

- Replay event logs to reconstruct order book

- Deterministic processing ensures consistent state

## 8.2 Checkpointing

Periodic snapshots of:

- order book state
- executed trades
- pending orders

## 8.3 Redundancy

Active-passive or active-active matching engine nodes.

---

# 9. Testing the Trading System

## 9.1 Unit Tests

- Matching scenarios
- Edge cases
- Partial fills
- Cancel logic

## 9.2 Latency Benchmarks

Measure:

- Gateway → Matching Engine latency
- Order processing time
- Market data publish time

## 9.3 Stress Testing

Simulate:

- Bursts of 1M orders
  - Network failures
  - Engine failover
- 

# 10. Business Impact of Low Latency (Bonus Section)

Latency has direct financial consequences.

## **10.1 Revenue Per Microsecond**

In market-making:

- Faster order placement = capturing spreads
- Faster cancellations = avoiding adverse selection
- Faster market data processing = better pricing decisions

A **100μs improvement** can translate to:

- Millions of dollars saved annually
- Lower slippage
- Better client execution quality

## **10.2 Cost–Benefit Analysis**

Investments in:

- Faster network cards
- Kernel bypass
- In-memory data grids

...are justified when:

- Increase in trading volume
- Reduction in failed trades
- Compliance improvements  
...exceed infrastructure cost.

## **10.3 Competitive Advantage**

A low-latency system allows:

- Better fill ratios
- Higher priority in matching
- Reduced risk exposure

This is why high-frequency trading firms and banks spend millions optimizing microseconds.

---

## **11. Conclusion & Future Scope**

This whitepaper presented a full technical design of a low-latency matching engine, detailing:

- Architecture and components
- Order book structures
- Matching algorithms
- Hot-path optimizations
- Fault tolerance and recovery
- Asynchronous persistence
- Business value of speed

**Future enhancements may include:**

- Hardware acceleration
- FPGA-based matching
- Advanced risk modeling
- Cross-venue optimization

---

## 12. References

*(You may include these — they look professional)*

1. “Matching Engine Design: Nasdaq Technical Briefs”
  2. “High Frequency Trading Systems Architecture,” ACM Queue
  3. “Kafka as a High Performance Distributed Log,” LinkedIn Engineering
  4. “Low Latency Networking Using DPDK,” Intel Documentation
  5. “Order Book Data Structures,” Quantitative Finance Journals
-